



GradCAM in PyTorch

400067-Prut




Why did our model make those predictions?


train set prediction: charmander



train set prediction: bulbasaur




test set prediction: charmander



Gradient-weighted Class Activation Mapping

Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization


3




(a) Original Image




(b) Guided Backprop 'Cat'




(c) Grad-CAM 'Cat'



(d) Guided Grad-CAM 'Cat'



(e) Occlusion map 'Cat'



(f) ResNet Grad-CAM 'Cat'




(g) Original Image




(h) Guided Backprop 'Dog'



(i) Grad-CAM 'Dog'



(j) Guided Grad-CAM 'Dog'



(k) Occlusion map 'Dog'

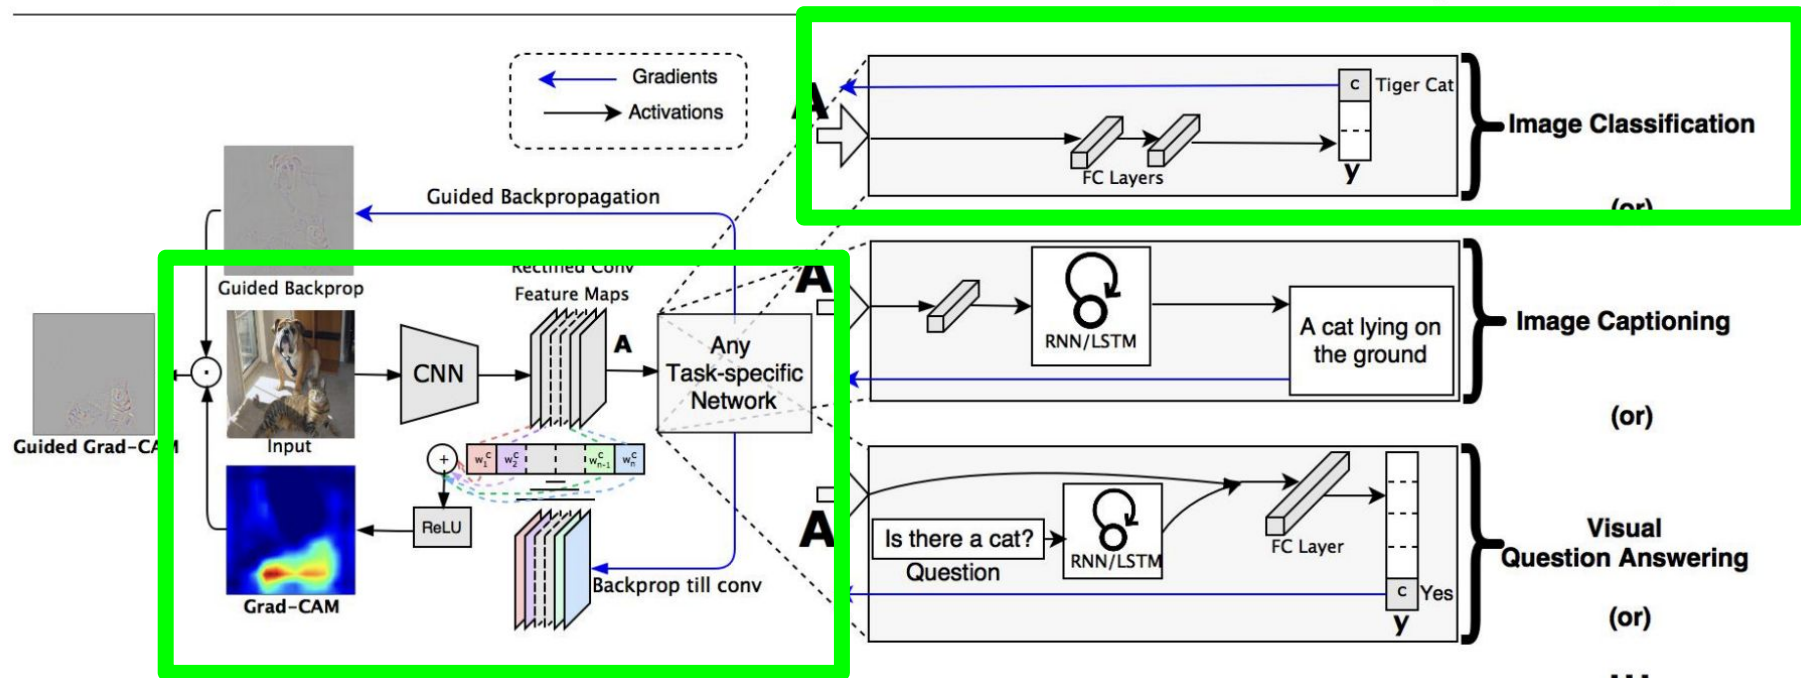


(l) ResNet Grad-CAM 'Dog'

GradCAM vs CAM

4

Ramprasaath R. Selvaraju et al.



```
!pip install grad-cam
```

```
from pytorch_grad_cam import GradCAM
```

<https://github.com/jacobgil/pytorch-grad-cam>



Code

<https://www.kaggle.com/code/prutsaowaaprut/gradcam-tutorial>



Dataset

<https://www.kaggle.com/datasets/prutsaowaaprut/charmander-or-bulbasaur-duck-invasion/data>



```
class CNNModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 32, 5)
        self.conv2 = nn.Conv2d(32, 64, 5)
        self.conv3 = nn.Conv2d(64, 128, 3)
        self.conv4 = nn.Conv2d(128, 256, 5)
        self.fc1 = nn.Linear(256, 2)
        self.pool = nn.MaxPool2d(2, 2)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = self.pool(F.relu(self.conv4(x)))
        x = F.adaptive_avg_pool2d(x, 1).reshape(x.shape[0], -1)
        x = self.fc1(x)
        return x
```

```
from pytorch_grad_cam import GradCAM
from pytorch_grad_cam.utils.model_targets import ClassifierOutputTarget
from pytorch_grad_cam.utils.image import show_cam_on_image
```

```
cam = GradCAM(model=model, target_layers=[model.conv4])
```

```
targets = [ClassifierOutputTarget(pred_class)]
```


```
grayscale_cam = cam(input_tensor=input_tensor.unsqueeze(0), targets=targets, aug_smooth=True)
```

```
grayscale_cam = grayscale_cam[0, :]
```


```
visualization = show_cam_on_image(torch_to_numpy(input_tensor), grayscale_cam, use_rgb=True)
```





train set prediction: charmander





test set prediction: charmander



test set prediction: charmander



Training set is biased!




Applications

- Object detection
- Image segmentation


Tuberculosis_classification_DenseNet121_GradCAM

35

Notebook Input Output Logs Comments (11)



<https://www.kaggle.com/code/sanphats/tuberculosis-classification-on-densenet121-gradcam#Grad-cam-evaluation>



Method**What it does**

| | |
|-----------------------------|--|
| GradCAM | Weight the 2D activations by the average gradient |
| HiResCAM | Like GradCAM but element-wise multiply the activations with the gradients; provably guaranteed faithfulness for certain models |
| GradCAMElementWise | Like GradCAM but element-wise multiply the activations with the gradients then apply a ReLU operation before summing |
| GradCAM++ | Like GradCAM but uses second order gradients |
| XGradCAM | Like GradCAM but scale the gradients by the normalized activations |
| AblationCAM | Zero out activations and measure how the output drops (this repository includes a fast batched implementation) |
| ScoreCAM | Perbute the image by the scaled activations and measure how the output drops |
| EigenCAM | Takes the first principle component of the 2D Activations (no class discrimination, but seems to give great results) |
| EigenGradCAM | Like EigenCAM but with class discrimination: First principle component of Activations*Grad. Looks like GradCAM, but cleaner |
| LayerCAM | Spatially weight the activations by positive gradients. Works better especially in lower layers |
| FullGrad | Computes the gradients of the biases from all over the network, and then sums them |
| Deep Feature Factorizations | Non Negative Matrix Factorization on the 2D activations |

Other Explainable AI (XAI) Methods

- LIME
- SHAP



References

Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2019). Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization. Retrieved from <https://arxiv.org/abs/1610.02391>

<https://www.analyticsvidhya.com/blog/2023/12/grad-cam-in-deep-learning/>

<https://github.com/jacobgil/pytorch-grad-cam>

<https://pytorch.org/vision/stable/index.html>

<https://debuggercafe.com/pytorch-imagefolder-for-training-cnn-models/>

<https://www.kaggle.com/code/sanphats/tuberculosis-classification-densenet121-gradcam#Grad-cam-evaluation>

<https://www.kaggle.com/competitions/rsna-breast-cancer-detection/discussion/372186>